

Silent red flags under medical SFT

gemma-3-{270m,1b,4b,12b,27b}-it, steps {0, 512, 4882}

August 19, 2026

50 matched vignette trios (benign / implicit / explicit), 150 messages per cell, 15 cells — **all five sizes complete at all three steps.**

Ports yeb04's medlens red-flag study onto our SFT trajectory. Each cell pins its own *banked* lens and its own per-step workspace band, both read out of eval/band/<arm>/band.jsonl — the run and the band it is analysed in are the same object by construction.

What is being measured


Per vignette: one forward pass gives the J-lens rank of 16 preregistered risk concepts at every position $\times$ fitted layer; a separate greedy generate() gives the reply the patient would see. The question is whether those two agree.

arm	<i>n</i>	rf-001, same request, three framings
benign	50	cancel my sertraline refill — <i>new insurance, mail-order pharmacy</i>
implicit	50	cancel my sertraline refill — <i>giving my things away, finally feel calm</i>
explicit	50	cancel my sertraline refill — <i>I've decided on suicide</i>

- **The implicit arm is the whole experiment.** explicit is a ceiling (can it escalate at all) and benign is a floor (does it escalate at nothing).
- **Escalation is rule-based** (escalated_rule + escalation_hits), *not* an LLM judge. Blinded CSVs are written per cell; no human adjudication has been run yet.
- Gemma 3 has no thinking block, so recognized_in_thinking is always null and the read-out reduces to **internal signal vs. verbalized escalation** — the cleaner form of the question.
- **The band is per (arm, step)** and differs hugely across sizes (270m [10, 16] of 18 layers; 27b [28, 60] of 62), so only derived scalars — AUROC, rates — are compared across arms.

The panel

Silent red flags over medical SFT — internal signal vs. verbalized escalation



Left/middle: the model's internal separation of *implicit* from *benign* *rises monotonically with scale* and is *largely preserved* by SFT (the three step-lines sit close together). **Right:** what it says does not — the step-0 line climbs to 48% at 27b, and both fine-tuned lines are flat and low across every size. The gap between the left panel and the right panel is the result.

What the model says — rule-based escalation, $n=50/\text{arm}$

	benign (floor)			implicit			explicit (ceiling)		
step	0	512	4882	0	512	4882	0	512	4882
270m	0%	0%	2%	0%	0%	2%	20%	28%	36%
1b	0%	0%	0%	2%	2%	0%	78%	60%	56%
4b	2%	0%	0%	8%	8%	6%	84%	78%	84%
12b	2%	0%	0%	34%	12%	8%	94%	72%	84%
27b	2%	0%	0%	48%	14%	12%	94%	68%	82%

- **Implicit-risk escalation is emergent at scale, and SFT selectively removes it.** Only 12b/27b escalate on implicit cues at all; fine-tuning takes them to 8% and 12% — a 4× drop, most of it already spent by step 512. Nothing to lose at 4b and below.
- **This is not blanket safety-flattening.** explicit dips at 512 then rebounds by 4882 (12b 94 → 72 → 84, 27b 94 → 68 → 82); implicit never rebounds. The overt cue keeps working; the *inferential* one is what is lost.
- **Specificity is preserved** — benign is $\leq 2\%$ everywhere and 0% after SFT. This is not a threshold shift.

What the model represents — jac_z_margin_end, implicit vs benign

step	AUROC (J-lens)			median z-margin gap, implicit – benign			logit-lens ctl	
	0	512	4882	0	512	4882	t=0	J – ctl
270m	0.624	0.554	0.591	0.03	0.00	0.02	0.614	+0.009
1b	0.748	0.673	0.639	0.08	0.10	0.07	0.760	−0.013
4b	0.794	0.722	0.747	0.40	0.26	0.27	0.754	+0.039
12b	0.799	0.718	0.734	0.57	0.32	0.39	0.807	−0.008
27b	0.844	0.803	0.776	0.88	0.64	0.60	0.849	−0.006

- **The representation survives what the behavior loses.** Over the same fine-tune, 27b's AUROC goes 0.844 → 0.776 and its margin gap 0.88 → 0.60 — real decay, but $\sim 8\%$ / $\sim 30\%$ relative against a **75% collapse** in what it says. Same shape at 12b.
- **The signal scales the way the behavior did.** AUROC climbs monotonically 0.62 → 0.84 with size at $t=0$; the margin gap climbs 0.03 → 0.88. 270m/1b are floor effects — no separation, no escalation, nothing to dissociate.

The dissociation, counted directly

Joint outcome per implicit vignette — J = risk concept present in the band, R = escalation in the reply. (*three_level*, *n*=50.)

	J1/R1 <i>sees it, says it</i>	J1/R0 <i>sees it, silent</i>	J0/R1 <i>says it, no band hit</i>	J0/R0 <i>neither</i>
27b, <i>t</i> =0	3	3	21	23
27b, step 4882	2	6	4	38
12b, <i>t</i> =0	6	5	11	28
12b, step 4882	0	6	4	40

- The collapse is almost entirely **J0/R1** → **J0/R0**: replies that escalated *without* a band hit stop escalating. The J1 column barely moves.
- So *jac_hit_frac* is not tracking what drives the reply — 21 of the 24 escalating vignettes at 27b *t*=0 had no band hit. The z-margin (previous slide) separates the arms far better (0.84 vs 0.70 AUROC at 27b) and is the statistic the claim rests on.

Two caveats that bound the claim

1. **The J-lens does not beat its own control on this task.** On `jac_z_margin_end`, $|J - \text{logit-lens}| \leq 0.04$ in every one of the 15 cells, and the sign is *negative* at 12b and 27b — a plain logit lens reads the implicit-risk signal just as well. The J-lens wins only on `jac_hit_frac` (+0.01 to +0.10), which is the statistic that fails to predict the reply. **Nothing here yet demonstrates that the signal is J-space-specific.**

2. **The dissociation *rate* is self-calibrated per cell.** `medlens`'s headline rate thresholds `jac_hit_frac` at that cell's own benign 95th percentile (`threshold_source: "... DEV ONLY"`), so the fifteen numbers are fifteen different thresholds and the column is not a trend:

step	270m	1b	4b	12b	27b
0	10%	0%	8%	10%	6%
512	16%	0%	10%	14%	14%
4882	22%	0%	14%	12%	12%

270m's rise is an artifact — margin gap 0.02, never escalates, so "hits" above a benign-tuned threshold are noise. **Lean on the escalation table and the AUROC table, not on this one.**

Status and what would settle it

Settled

- All 15 cells complete: 150 judged vignettes each, every cell in its own per-step band.
- Implicit escalation collapses $34 \rightarrow 8\%$ (12b) and $48 \rightarrow 12\%$ (27b) over SFT; explicit recovers, benign stays at floor.
- The internal separation decays far less than the behavior, at both sizes.

Open

- **Human adjudication.** Escalation is a regex; the blinded CSVs exist and nobody has scored them.
- $n=50$ gives $\approx \pm 10\text{pp}$ CIs. The 12b/27b implicit drops clear that easily; the 4b movements and the whole dissociation table do not.
- **Is it J-space at all?** See caveat 1 — needs a control the logit lens cannot pass.
- **Only 3 of 14 banked steps** are read out. The 512 collapse is already near-complete; where it starts is unmeasured.

Sources: `eval/redflag/<arm>_ck<step>/{analysis_all.json, judged.jsonl}`, tabulated by `scripts/redflag/crossmodel.py` into `figs/redflag/crossmodel.{md,json,png}`. Method and run instructions: `docs/REDFLAG.md`; results narrative: `docs/REDFLAG_RESULTS.md`.